

Proposed Solution

1. Proposed Architecture Overview

The proposed solution wraps the predictive scikit-learn models inside a state-of-the-art Python Flask MVC architecture. We incorporate a multi-step user onboarding interface, full-page loading indicators, dynamic theme preference persistence via localStorage, and a print-ready PDF certificate generator complete with automated QR validation links.

2. Expected Metrics

- Pipeline latency: <15ms per request.
- Scoring precision: >85% balanced accuracy.
- Serverless footprint: <245 MB bundle size.